Image Captioning

Basic Definitions

Computer Vision

Natural Language Processing

Image Captioning

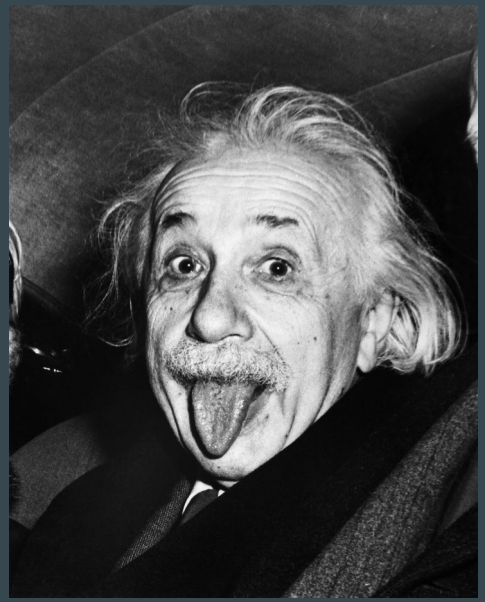

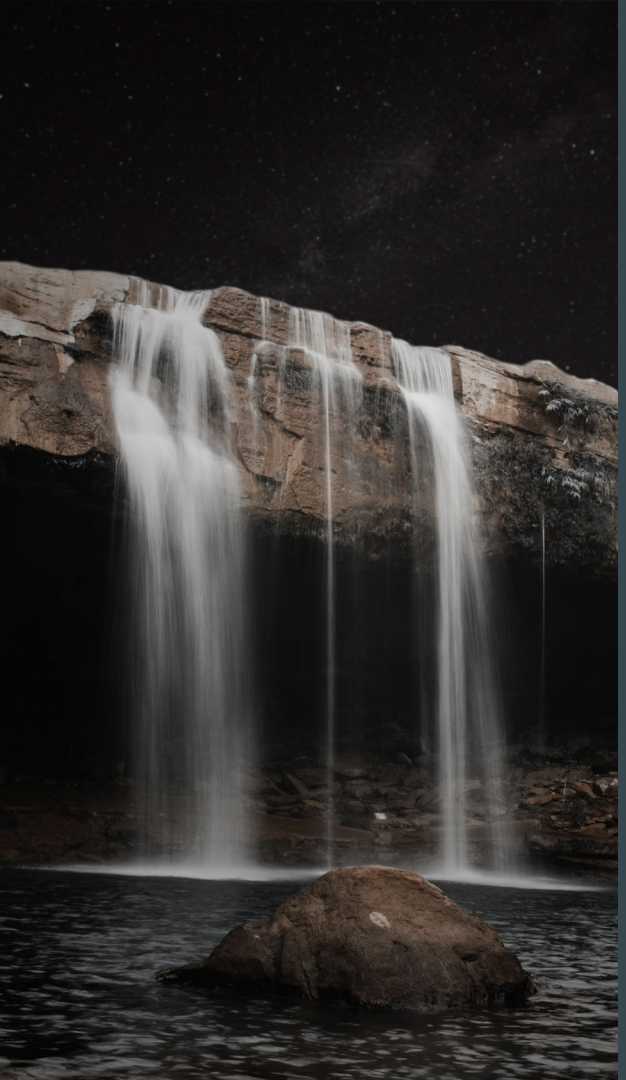

::Motivations::

Google Image Search Results

Reverse searching images using their captions

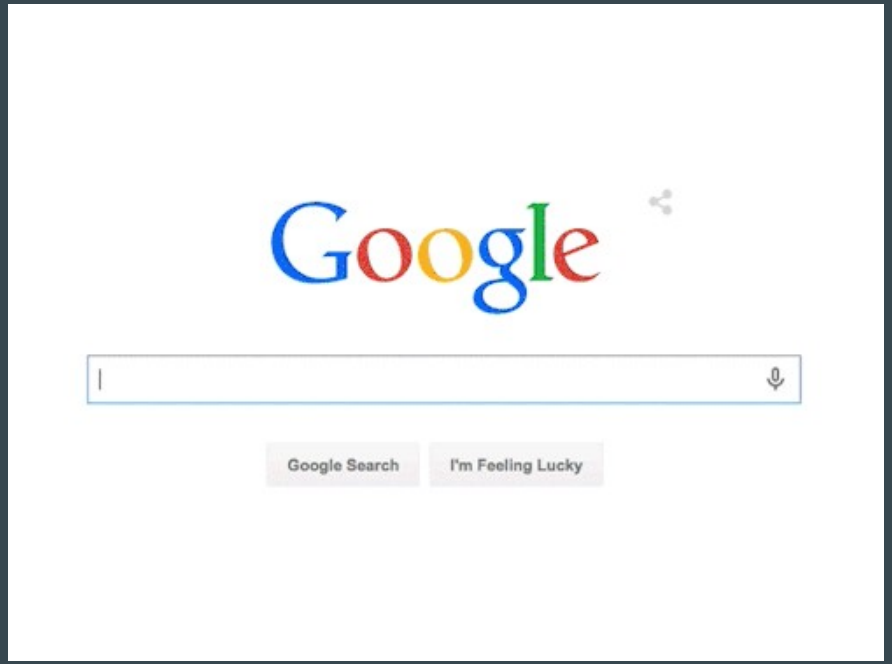

Aid to Blind

Describing
Surroundings

Self Driving Vehicles

Logging Journeys

Flickr8K

More than 8000 images with 5 captions per image

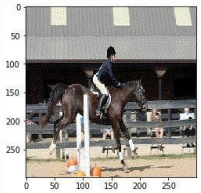

['startseq a jockey on a black horse jumps over a hurdle endseq',
'startseq an equestrian and a horse are jumping over an obstacle endseq',
'startseq a person wearing a navy jacket and black hat jumping over a small partition on a horse endseq',
'startseq a show jumper is making a brown horse jump over a white fence endseq',
'startseq a woman on a horse jumps an obstacle endseq']

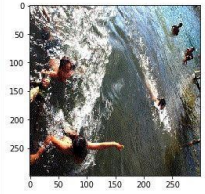

['startseq a bunch of people swimming in water endseq',
'startseq a group of children in the ocean endseq',
'startseq a group of youngsters swim in lake water endseq',
'startseq many children are playing and swimming in the water endseq',
'startseq several people swim in a body of water endseq']

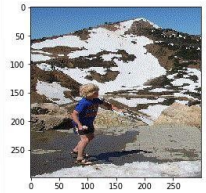

['startseq a blond hair boy in short short sleeve shirt and sandals in overlooking a snowcapped mountain endseq',
'startseq a boy in a blue shirt is standing at the foot of a hill with a snowball in his hand endseq',
'startseq a boy in a t shirt and shorts is holding a snowball and facing a snowy mountain endseq',
'startseq a boy preparing to throw a snowball endseq',
'startseq a child in shorts throws a snowball at a mountain endseq']

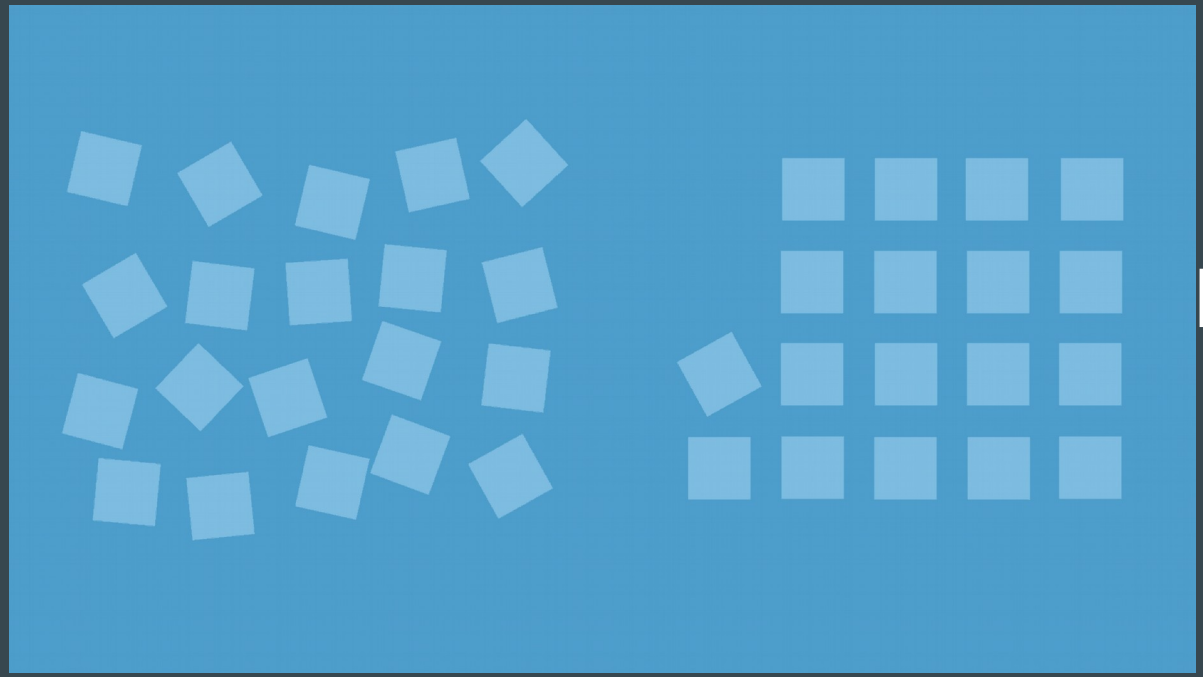

Preprocessing

Processing captions

- Turned into lowercase.
- Numbers are removed.
- Start/End tags are added

Tokenization

- Assigning numbers to unique words.

START a little girl in pink climbs a rope bridge at the park . END

tensor([[3,   2, 41, 20,   5, 91, 252,   2, 212, 333, 23,   6, 119,   4]])

Architecture

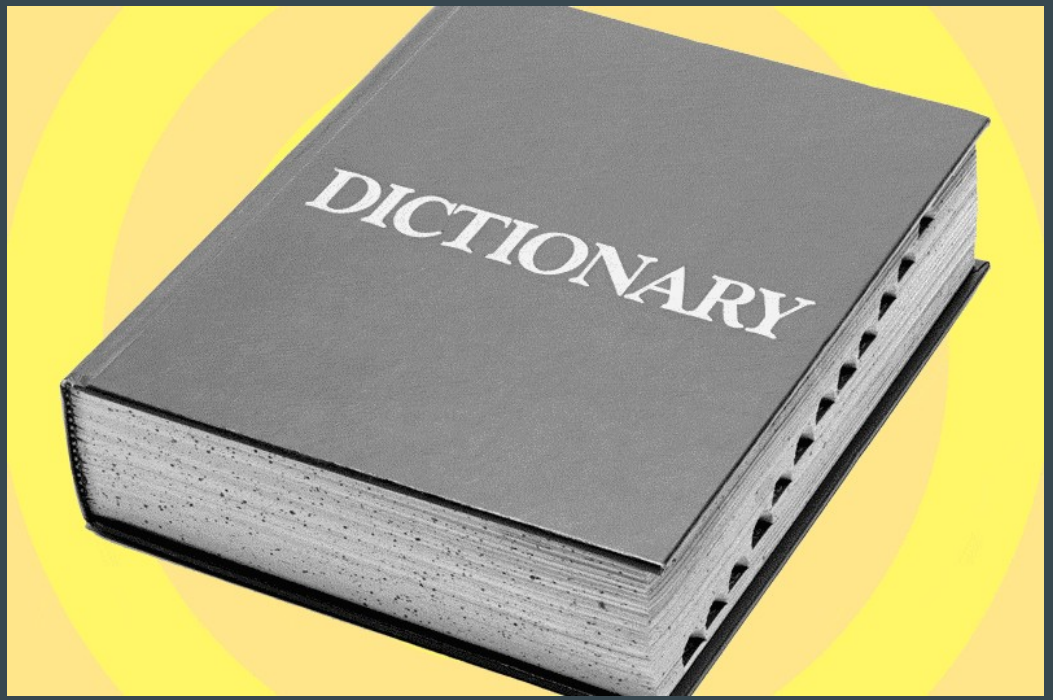

AutoEncoder

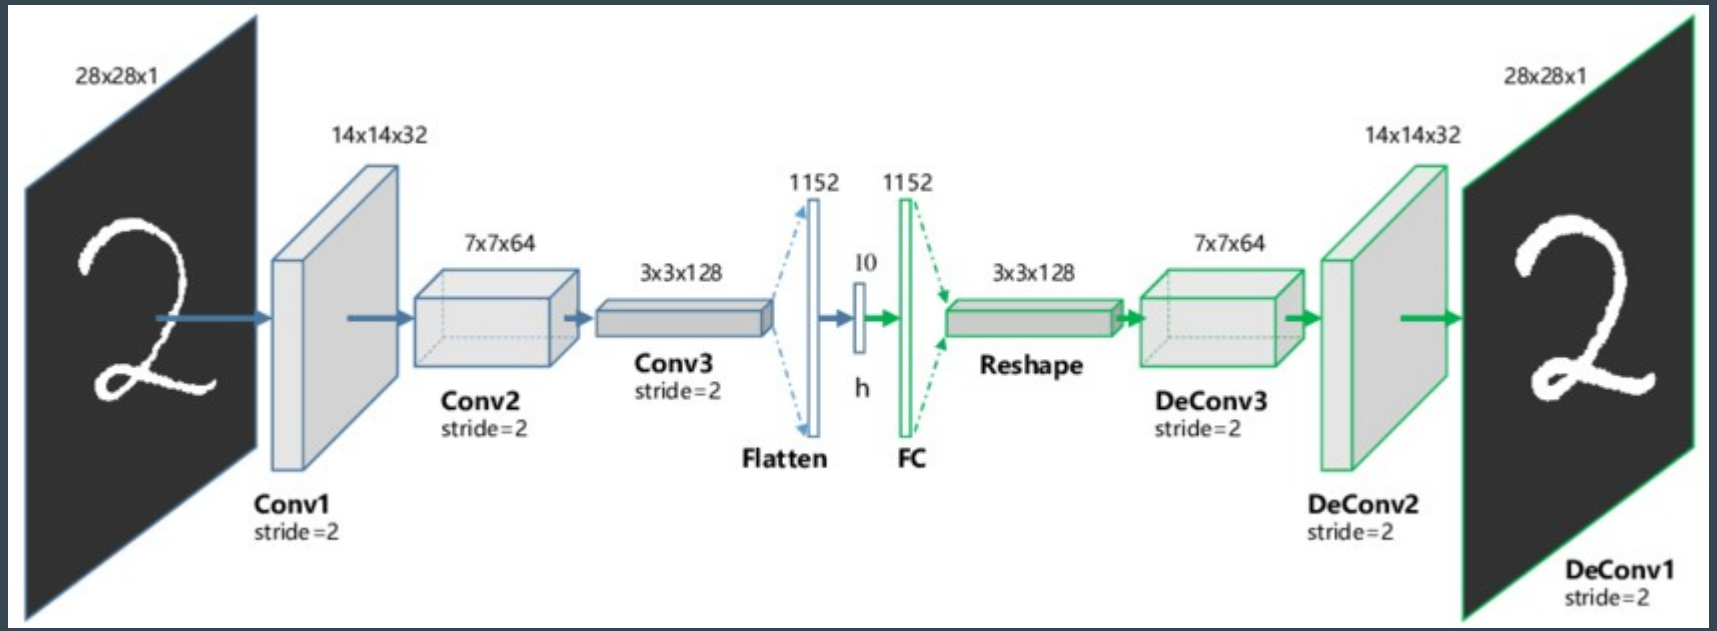

Encoder - DenseNet

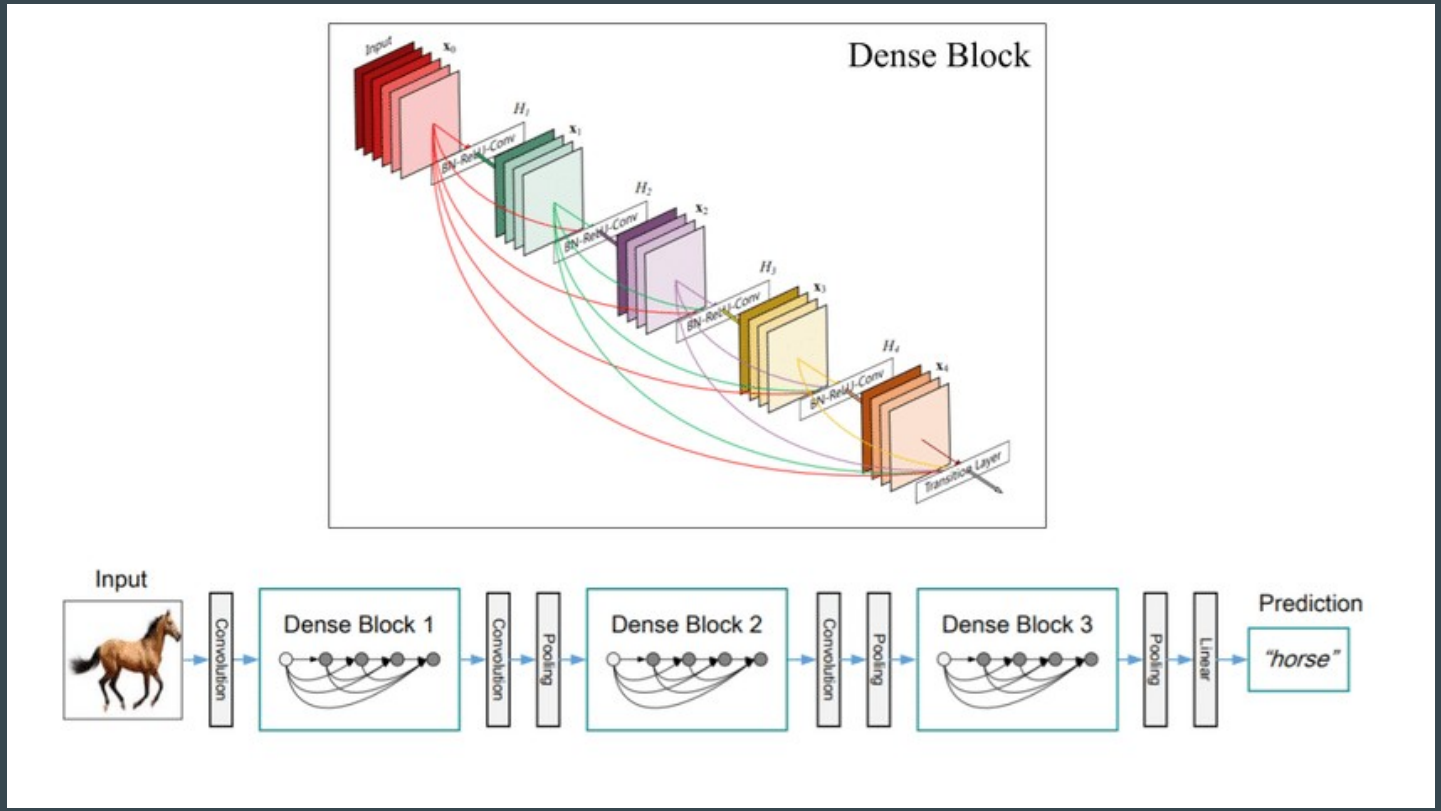

Attention

Decoder - GRU with Attention

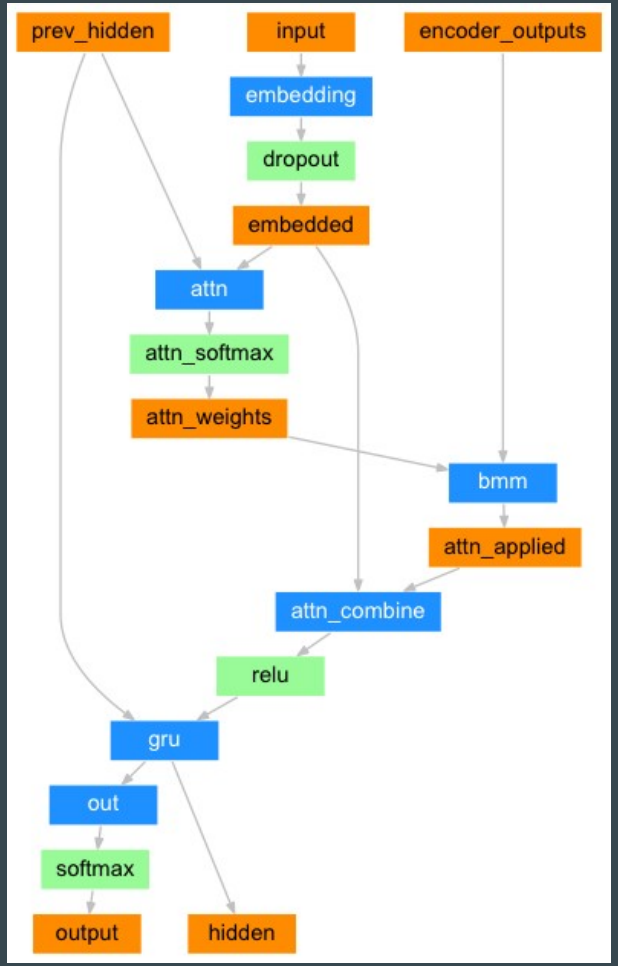

<code/>

Some results from our Model

[0/75.0%] loss: 4.237801445855035 time: 23.803110122680664 sec
a young dog runs to catch a ball END -->
START a young boy running with a boogie board into the water END

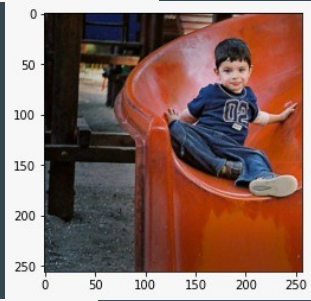

[1/90.0%]loss:2.369701385498047 time:0.0034105579058329263mins
a dog boy is a black dress and down a street slide END -->
START a young boy wearing a blue outfit sliding down a red slide . END

[0/25.0%] loss: 4.34010021503155 time: 23.559677839279175 sec
a dogs run jumping a a a a beach END a a END -->
START three dogs run through the water near the rocks and make splashes . END

Conclusion

- Results may differ
- Accuracy ~
  - Larger Dataset
  - Model architecture
  - Hyperparameter Tuning
  - Using Cross Validation